

Niko Kirste

Faculty of Computer Science

Real-time Lighting for Mixed Reality Applications

Final Presentation, INF-PM-FPA

Dresden, 16.10.2024

Table of Contents

- **Motivation**
- **Research Question**
- **Existing Methods & Related Work**
- **Implementation Process**
- **Results**
- **Future Work**
- **Conclusion**

Motivation

Lighting in traditional VR Applications:

- Closed environment
- No interaction with real world lighting

→ static precomputed lighting
→ low overhead



Lighting in MR Applications:

- Ongoing interaction with real world lighting
- Constantly changing
- Immersion depends on correct lighting

→ need to adapt



Research Question

**How can we get lighting information from the physical world
to the MR application in real-time?**

- + **Accurate colors**
- + **Low overhead on the HMD**
- + **Preferably multi-user enabled**

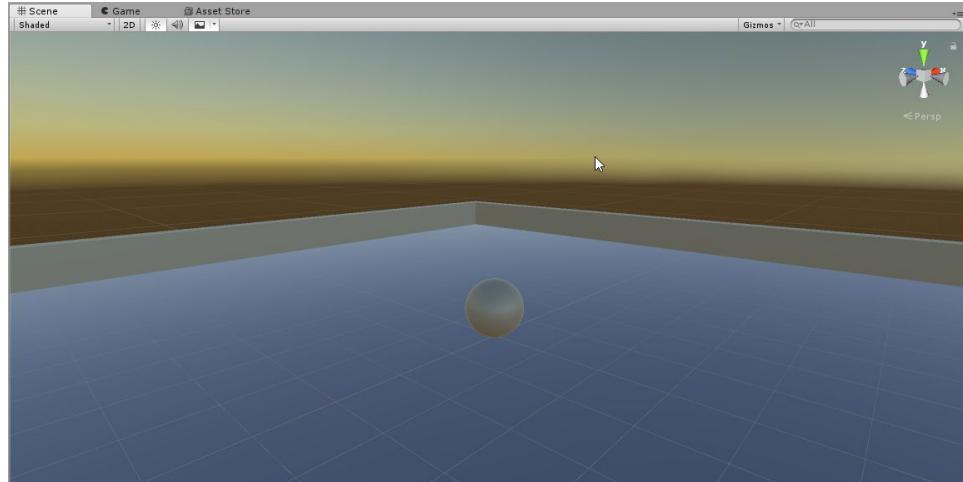
Existing Methods

Skybox:

- **Static or dynamic 360° texture**
- **Projected into infinity**
- **Contributes to single ambient light probe**

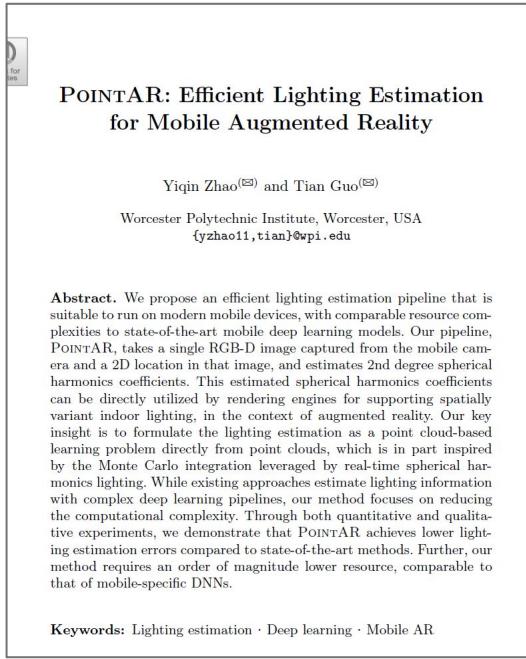
→ **no spatial gradient**

→ **same lighting for every object**



Related Work

Light Probe generation



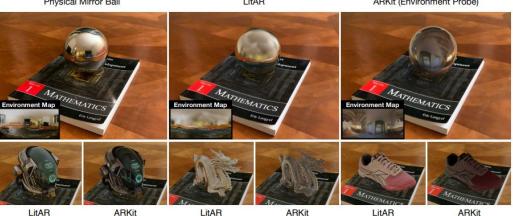
POINTAR: Efficient Lighting Estimation for Mobile Augmented Reality

Yiqin Zhao^(✉) and Tian Guo^(✉)

Worcester Polytechnic Institute, Worcester, USA
yzhao11,tian}@wpi.edu

Abstract. We propose an efficient lighting estimation pipeline that is suitable to run on modern mobile devices, with comparable resource complexities to state-of-the-art mobile deep learning models. Our pipeline, POINTAR, takes a single RGB-D image captured from the mobile camera and a 2D location in that image, and estimates 2nd degree spherical harmonics coefficients. This estimated spherical harmonics coefficients can be directly utilized by rendering engines for supporting spatially variant indoor lighting, in the context of augmented reality. Our key insight is to formulate the lighting estimation as a point cloud-based learning problem directly from point clouds, which is in part inspired by the Monte Carlo integration leveraged by real-time spherical harmonics lighting. While existing approaches estimate lighting information with complex deep learning pipelines, our method focuses on reducing the computational complexity. Through both quantitative and qualitative experiments, we demonstrate that POINTAR achieves lower lighting estimation errors compared to state-of-the-art methods. Further, our method requires an order of magnitude lower resource, comparable to that of mobile-specific DNNs.

Keywords: Lighting estimation · Deep learning · Mobile AR



LITAR: Visually Coherent Lighting for Mobile Augmented Reality

YIQIN ZHAO, Worcester Polytechnic Institute, USA
CHONGYANG MA, Kuashou Technology, China
HAIBIN HUANG, Kuashou Technology, China
TIAN GUO, Worcester Polytechnic Institute, USA

Physical Mirror Ball LITAR ARKit (Environment Probe)

LITAR ARKit LITAR ARKit LITAR ARKit

Figure 1. Rendered objects with LITAR compared to a physical mirror ball and with ARKit [34]. Row 1: LITAR produces visually-coherent rendering while ARKit fails; the bottom part of the ARKit-rendered ball does not faithfully reflect the book cover. The top part of the LITAR-rendered ball has an intentionally gradient blurring effect for quality-performance trade-offs (see §4.2). Row 2: LITAR achieves more realistic and visually coherent rendering effects than ARKit for objects with different materials.



Accidental Light Probes

Hong-Xing Yu¹ Samir Agarwala¹ Charles Herrmann² Richard Szeliski² Noah Snavely²
Jiajun Wu¹ Deqing Sun²

¹Stanford University ²Google Research

Abstract

Recovering lighting in a scene from a single image is a fundamental problem in computer vision. While a mirror ball light probe can capture omnidirectional lighting, light probes are generally unavailable in everyday images. In this work, we study recovering lighting from accidental light probes (ALPs)—common shiny objects like Coke cans, which often accidentally appear in daily scenes. We propose a physically-based approach to model ALPs and estimate lighting from their appearances in single images. The main idea is to model the appearance of ALPs by photogrammetrically principled shading and to invert this process via differentiable rendering to recover incidental illumination. We demonstrate that we can put an ALP into a scene to allow high-fidelity lighting estimation. Our model can also recover lighting for existing images that happen to contain an ALP.

I'd rather be Shiny. — Tamatoa from Moana, 2016

1. Introduction

In general, recovering scene illumination from a single view is fundamental for many computer vision applications such as virtual object insertion [1], relighting [43], and photorealistic data augmentation [51]. Yet, it remains an open problem primarily due to its highly ill-posed nature. Images are formed through a complex interaction between geometry, material, and lighting [21], and without precise prior knowledge of a scene's geometry or materials, lighting estimation is extremely under-constrained. For example, scenes that consist primarily of matte materials reveal little information about lighting, since diffuse surfaces behave like low-pass filters on lighting during the shading process [38], eliminating high-frequency lighting information. To compensate for

Related Work

Neural Network based Lighting

SCIENCE CHINA
Information Sciences

CrossMark
4-8 hr to update

• RESEARCH PAPER •

February 2021, Vol. 64, 64–122101-1–122101-15
<https://doi.org/10.1007/s11432-020-9242-5>

Neural compositing for real-time augmented reality rendering in low-frequency lighting environments

Shengjie MA^{1†}, Qian SHEN^{1†}, Qiming HOU^{1,2}, Zhong REN^{1,2} & Kun ZHOU^{1,2*}

¹*State Key Lab of CAD&CG, Zhejiang University, Hangzhou 310058, China;
²*ZJU-Pace/Unity Joint Lab of Intelligent Graphics, Hangzhou 310015, China**

Received 13 June 2020/Revised 13 July

Abstract We present neural compositing which uses convolutional neural networks to graph to emulate shadow and reflection effect information from the photograph using new color, shadow and reflection layers by applying shadow and reflection layers to the original image. We assume low-frequency lighting and layer rendering, which makes the whole pipe with synthetic scenes. Working on a single image with spatially-varying material and camera at real-time frame rates.

Keywords augmented reality, neural network

Citation Ma S J, Shen Q, Hou Q M, et al. A frequency lighting environments. Sci China Inf

Learning Lightprobes for Mixed Reality Illumination

David Mandl¹* Keang Moo Yoo² Peter Mohr¹ Pascal Fua³ Vincent Lepetit¹ Deter Schmalzgug¹ Denis Kalokin¹

¹Graz University of Technology ²EPFL ³University of Bordeaux

Figure 1: Coherent illumination. The real scene consists of an action figure of The Hulk. To demonstrate the result of our system, we place a 3D scan next to the real action figure and display it using Mixed Reality: (left) The 3D reconstruction is rendered without real world light estimation. (right) Our system estimates the current lighting from a single input image. The estimated lighting is used to illuminate the 3D reconstruction. Note that we only register the real world lighting and do not consider any camera effects such as exposure or blur.

1 Introduction

The availability of commercial AR (augmented reality) development tools to render virtual objects technologies estimate camera poses and locations making the rendered objects appear real is

A key challenge is the interaction between spatially-varying shadow and produce reflectance crucial to realism, yet difficult to simulate and material information of the real scene [1].

ABSTRACT

This paper presents the first quantitative pipeline for MR that based on the first quantitative pipeline using convolutional neural network (CNN). For quick adaptation of the system, we train the CNNs using prior synthetic data. The proposed pipeline is simple, fast, and accurate. As the pipeline is efficient and efficient, we propose to use the light estimation results from multiple CNN instances and show an approach for caching estimation results. In addition, we propose to use the pipeline multiple strategies for the CNN training. Experimental results show that the proposed pipeline can achieve better performance than the previous work.

proaches rely on active light probes, e.g., fishey cameras strategically placed in the scene [10], or passive light probes, like chessboard spheres, or other reference objects with known shape and reflectance [4]. An obvious major downside of having a light probe is that it is not always feasible to have one in every scene. Another approach [8] demonstrates a probe-less method that estimates the incident light in real time using a depth sensor and a camera. This approach requires artificial lightprobes, but suffers from high computational demands [8].

We propose *learned lightprobes* as an alternative to the probeless method of Graber et al. [8]. Like the probeless method, it preserves user experience, while still retaining the computational benefits of

Distributed Lighting

A Distributed, Decoupled System for Losslessly Streaming Dynamic Light Probes to Thin Clients

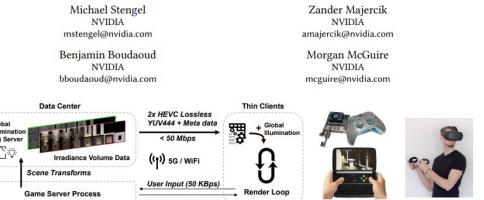


Figure 1. Light probe streaming architecture. Our system streams irradiance volume data from a cloud server to thin clients for mobile gaming and AR/VR. The streamed data enables dynamic high quality ray traced indirect lighting on the clients at low computational cost. We use efficient encoding and decoding hardware features to achieve high compression rates with low latency. Thus a single server can update thousands of light probes per seconds across multiple connected thin clients, amortizing rendering costs.

our split-rendering solution decouples remote computation from local rendering and so does not limit local display update rate or display resolution.

1 Introduction

Motivation. Today's high performance graphics systems approach cinematic rendering quality in real time using ray traced global illumination, with interaction latency measured in milliseconds. For applications like virtual and augmented reality (VR/AR), the need for low latency and high-fidelity rendering is augmented by a need for a small form factor and low thermal requirements to facilitate user comfort.

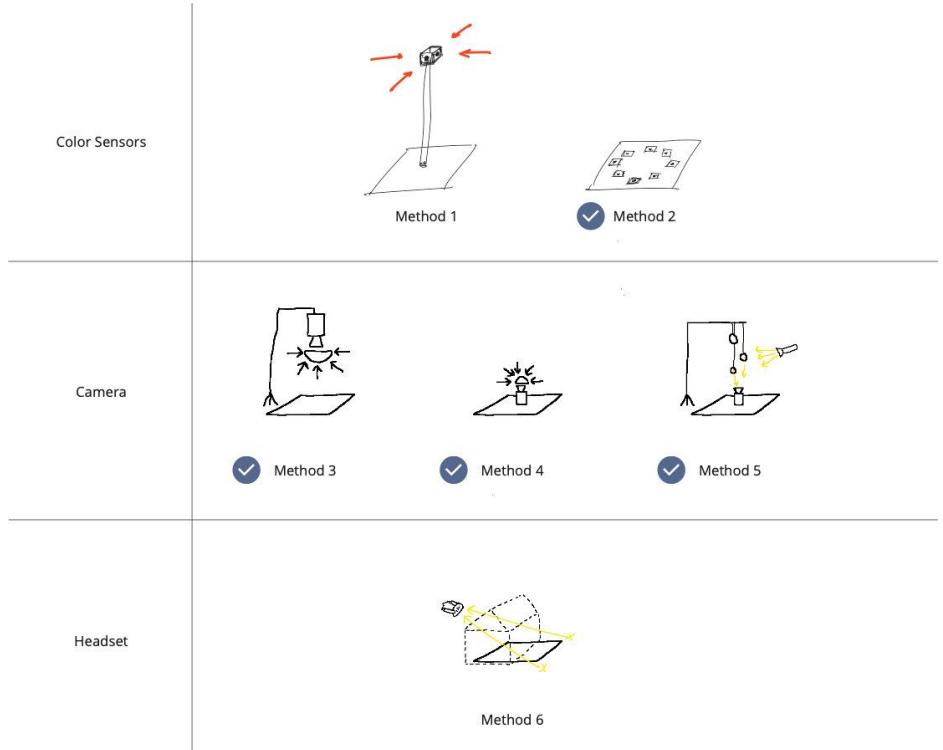
constrain high-fidelity rendering at low latency. This constraint becomes more severe as rendering requirements increase and desired form factor and thermal budget decrease, making it impossible to render high quality graphics directly on AR/VR/mobile devices.

Distributed Graphics. The state of the art for accelerat-

Distributed Graphics. The state of the art for accelerat-

Initial Ideas

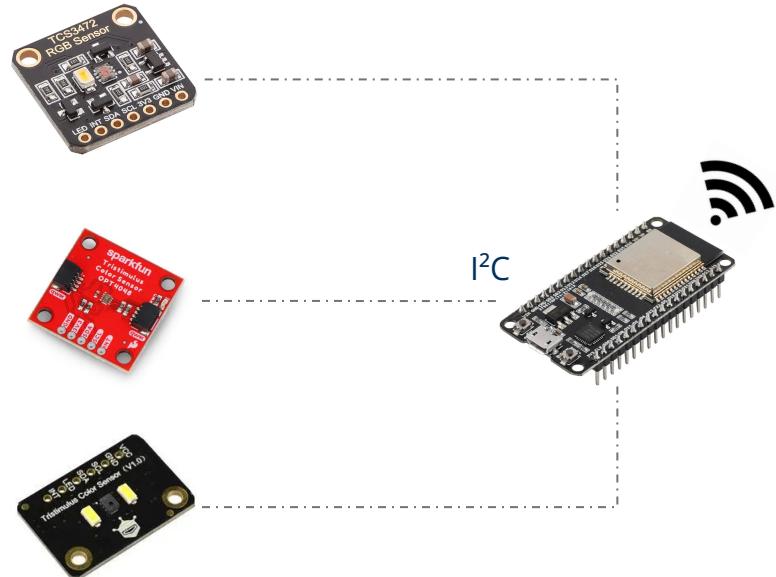
- Come up with different methods of capturing lighting information
- Implement and compare some methods
 - Feasibility
 - Visual Appearance
 - Cost
 - Complexity
- → Focusing on small scale / local methods for single objects



Color Sensor Methods

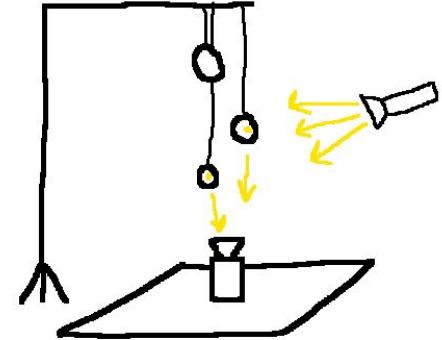
Hardware:

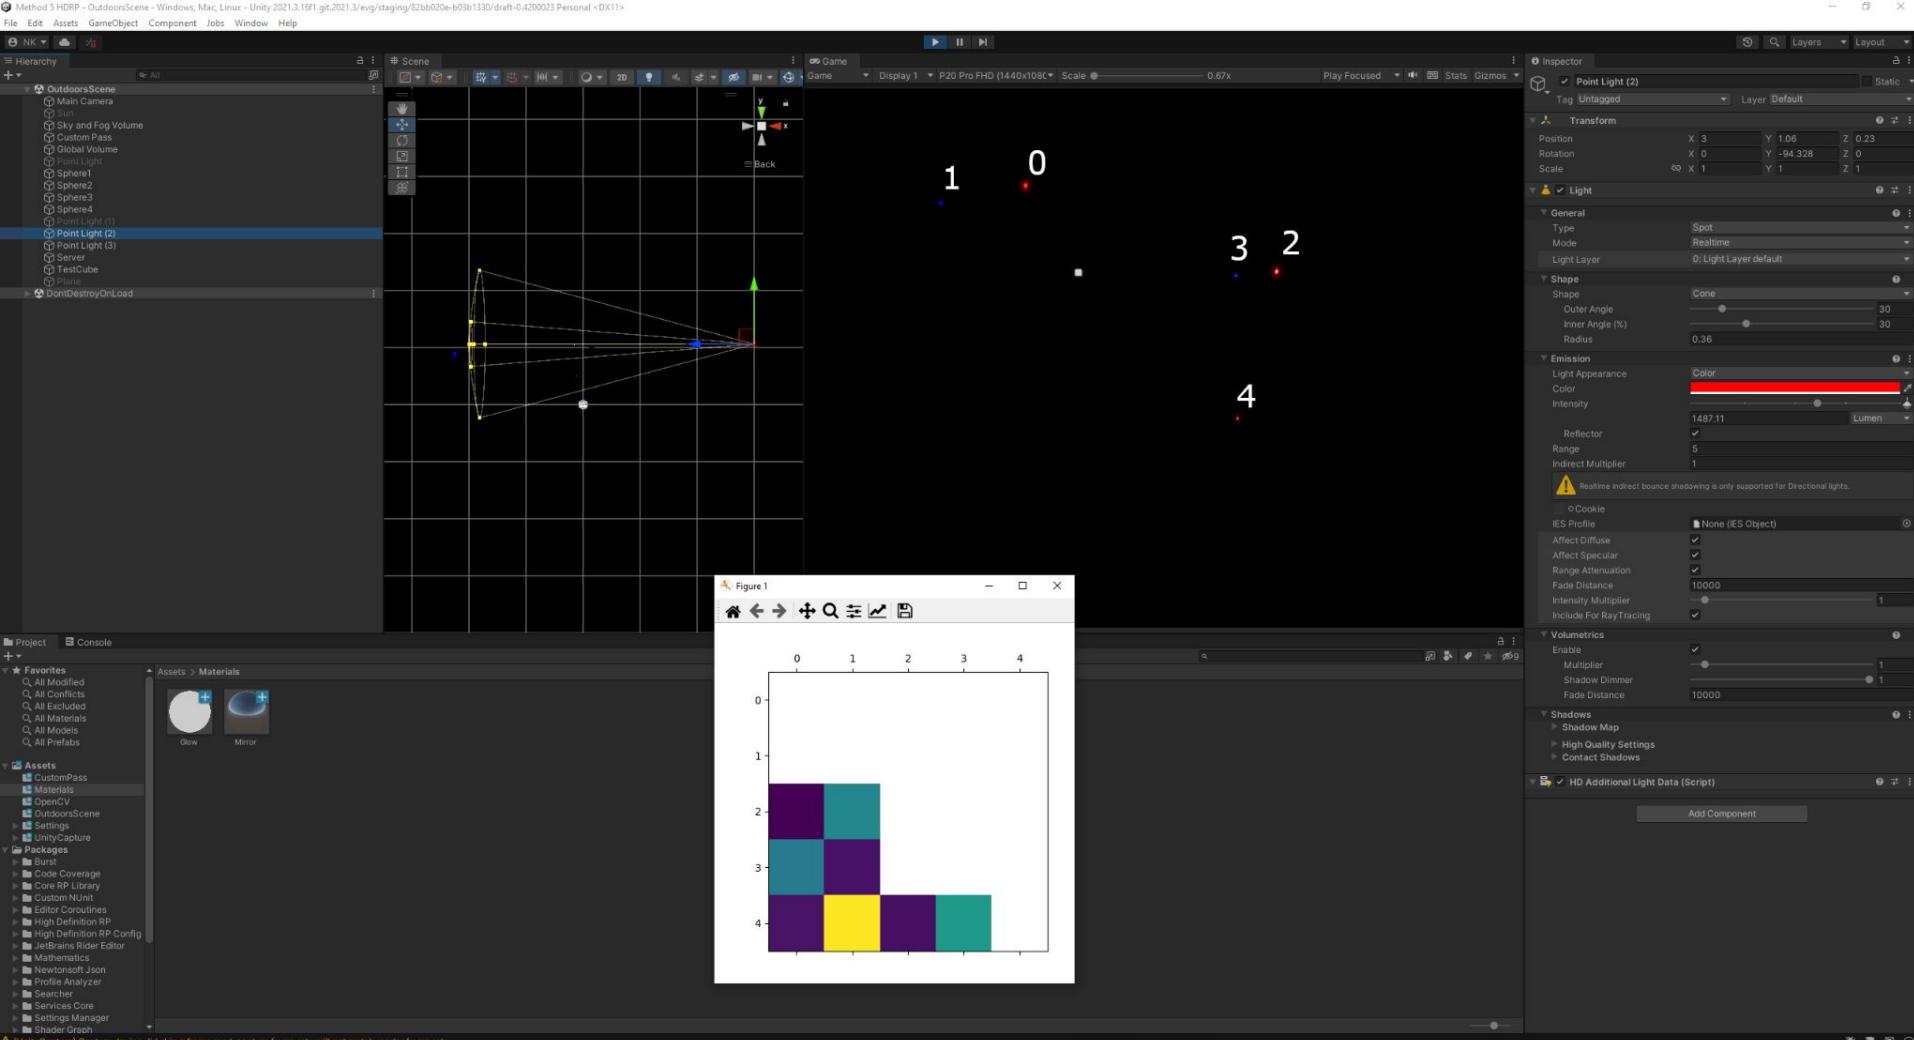
- the TCS34725:
 - RGBW sensor
 - I²C communication
 - cheap (~5€ per sensor, ~2€ from AliExpress)
- the OPT4048DTSR:
 - XYZ sensor
-> more accurate color sensing + light intensity in lux + color temperature
 - I²C communication
 - more expensive (~10€ per sensor)
- the SEN0403:
 - XYZ sensor (like OPT4048DTSR)
 - I²C communication
 - more expensive (~13€ per sensor)



Single Camera, multiple Spheres

- Calculate 3D position of light source
- Create virtual light with same properties at same position
- Affordable + Precise
 - Standard camera / smartphone
 - Some mirror spheres (e.g. christmas ornaments)





Hierarchy Scene Game Display 1 P20 Pro FHD (1440x1080) Scale 0.48x Play Focused Stats Gizmos

Inspector Server Static Tag Untagged Layer Default

Transform Posi X -0.15 Y 0.5 Z -0.15
Rota X 0 Y 0 Z 0
Scale X 1 Y 1 Z 1

IPC Server (Script)
Script: IPCServer
Transform: TestCube

Udp Receiver (Script)
Script: UdpReceiver
Positions: 10
Colors: 10
Light Objects: 10

Element: Light0 (Light)
Light1 (Light)
Light2 (Light)
Light3 (Light)
Light4 (Light)
Light5 (Light)
Light6 (Light)
Light7 (Light)
Light8 (Light)
Light9 (Light)

Project Console

Assets > OpenCV

brightness... detect_ligh... frame hdr_test image_line... intersec_t... iso_15 iso_15
iso_100_cu... light_analy... light_analy... light_match... optimize_p... ricoh_test test_iso_50 test_is

Distance matrix:
[[inf inf inf
[0.01266754 inf inf
[0.03174976 0.02751731 inf]]
2.599917674855853, -0.16033488758382206, 0.45762234431000776, 1, 0, 134
Light pairs indices:
[{0, 2, 3}]
Light pairs mean location:
[array([2.649363 , -0.18021096, 0.45948326])]
Mean Color:
[array([1, 0, 129])]
Distance matrix:
[[inf inf inf inf
[inf inf inf inf
[0.01708555 0.60531781 inf inf
[0.06643486 2.06871491 0.04046208 inf]]
2.6493629965783025, -0.1802109564789324, 0.4594832597486164, 1, 0, 129

! input string was not in a correct format.



Suchen



Improved Idea

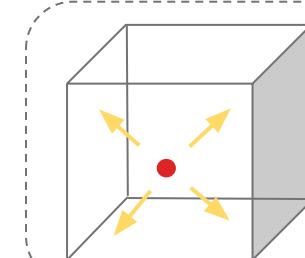
- **Use 360° camera instead of fish eye lens**
- **Use model of the room to project image onto**
- **Use emissive material to cast light into the scene**
- **Use light probes to light virtual objects**

System Architecture

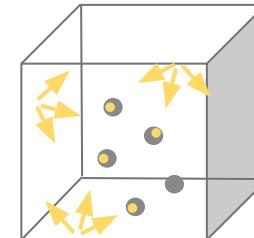
1. Capture Environment



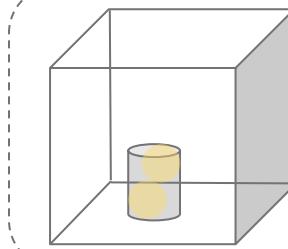
2. Project onto Walls



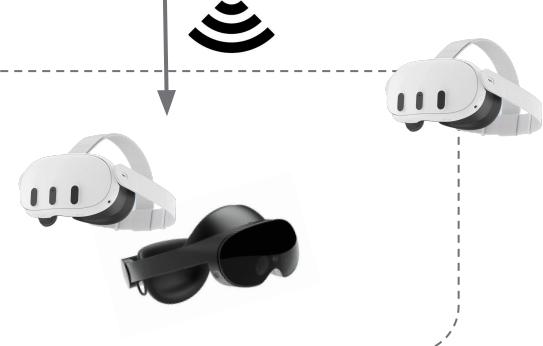
3. Calculate Light Probes



5. Light Virtual Objects

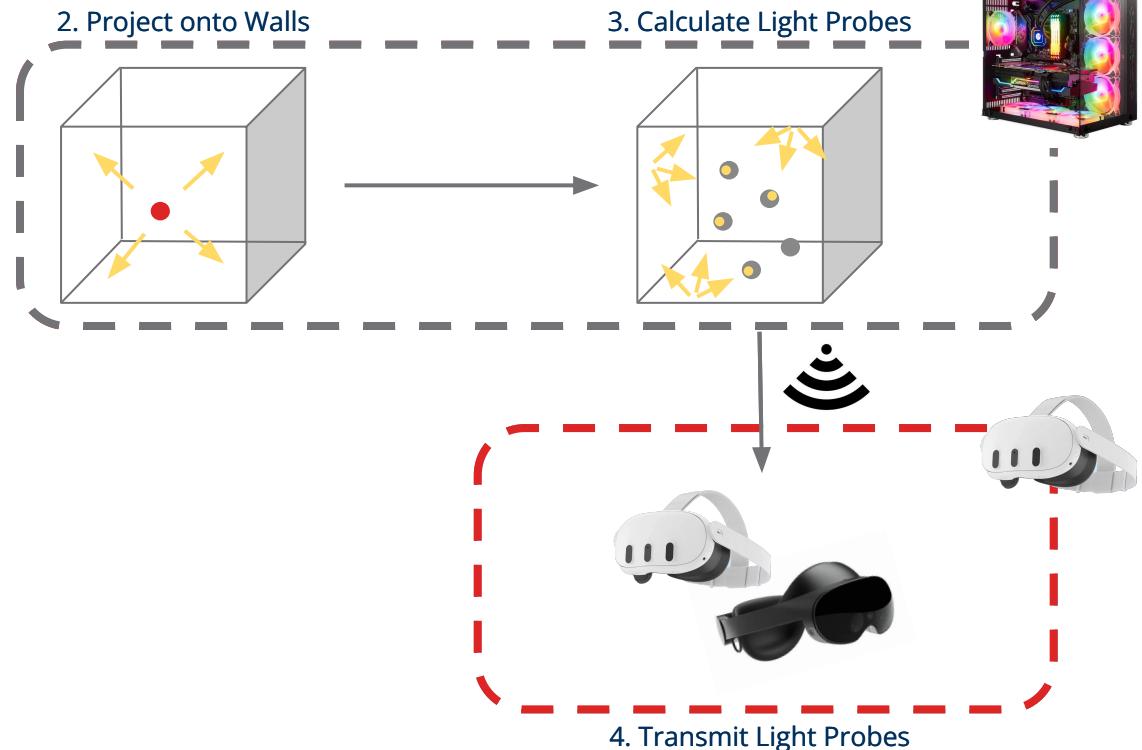


4. Transmit Light Probes



Key Features

1. Model based Lighting using emissive Materials

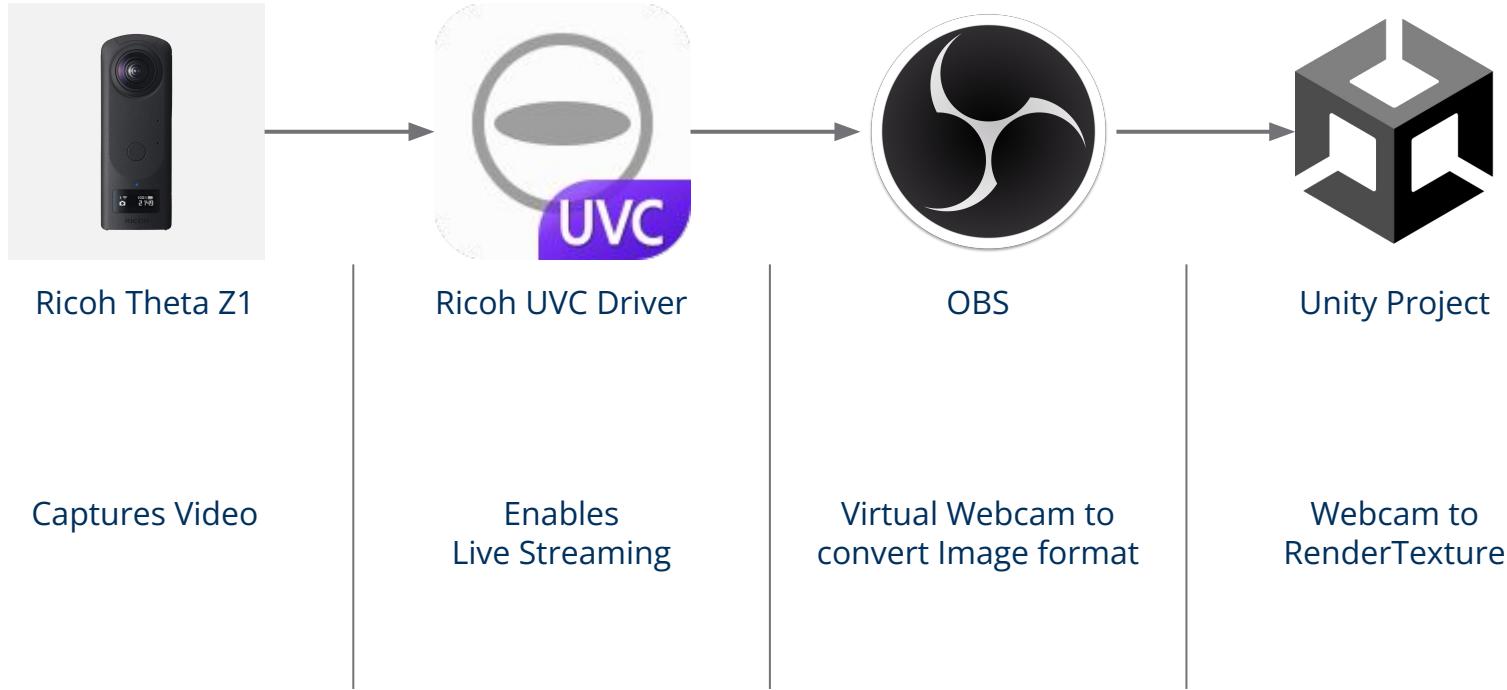


2. Streaming of Light Probes to multiple mobile Clients

1. Capture the Environment

	Ricoh Theta Z1		Vuze XR	
Lenses	2x 180° x 180°		2x 180° x 180°	
Video Resolution	Up to 4K @ 30FPS		Up to 5.7K @ 30FPS	
HDR support	For photos		None	
Live Streaming	Over USB or per App		Per App	
Software Availability	Over their Website		Discontinued	

1. Capture the Environment



1. Capture the Environment

```
5 public class WebcamToRenderTexture : MonoBehaviour
6 {
7     public RenderTexture outputTexture;
8     public Texture image;
9
10    public string cameraDeviceName = "OBS";
11    private WebCamTexture webCamTexture;
12
13    public bool enableWebcam = false;
14    public bool enableImage = false;
15
16    void Start()
17    {
18        if (!enableWebcam) return;
19
20        WebCamDevice[] devices = WebCamTexture.devices;
21        WebCamDevice device = devices[0];
22
23        // Prints available webcams to the console
24        for (int i = 0; i < devices.Length; i++)
25        {
26            print("Webcam available: " + devices[i].name);
27            if (devices[i].name.Contains(cameraDeviceName))
28            {
29                device = devices[i];
30            }
31        }
32
33        webCamTexture = new WebCamTexture(device.name);
34        webCamTexture.Play();
35    }
36
37    void Update()
38    {
39        if (enableWebcam)
40        {
41            Graphics.Blit(webCamTexture, outputTexture);
42        }
43        if (enableImage)
44        {
45            Graphics.Blit(image, outputTexture);
46        }
47    }
}
```

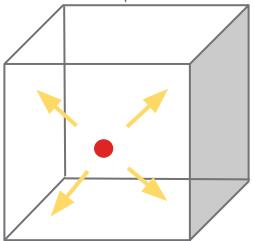
- Selectable webcam device
- Use static images for testing



2. Projection onto Walls



UV-Mapping



```
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111

void CreateUvMap(GameObject gameObject)
{
    if (gameObject.GetComponent<MeshFilter>() == null) return;

    Mesh mesh = gameObject.GetComponent<MeshFilter>().mesh;

    // merge duplicate vertices from previous mappings
    NormalizeMesh(mesh);

    // calculate the UVs without post-correction
    CalculateUvs(mesh, gameObject);

    // correct the edges that span over the edge of the UV map
    // using a vertex duplication strategy
    FixUVSeam(mesh, gameObject);
}

// Calculate the UV coordinates of every vertex depending on its world position
// and the center of projection
public void CalculateUvs(Mesh mesh, GameObject gameObject)
{
    Vector3[] vertices = mesh.vertices;
    Vector2[] uvs = new Vector2[vertices.Length];

    for (int i = 0; i < uvs.Length; i++)
    {
        Vector3 worldPosition = gameObject.transform.TransformPoint(vertices[i]);
        Vector3 sphereCenter = ProjectionCenter.transform.position;

        Vector2 uv = GetSphereUV(worldPosition, sphereCenter);
        uv = new Vector2(uv.x, 1.0f - uv.y);
        uvs[i] = uv;
    }

    mesh.uv = uvs;
}

public Vector2 GetSphereUV(Vector3 point, Vector3 sphereCenter)
{
    Vector3 normalizedPoint = (point - sphereCenter).normalized;
    normalizedPoint = Quaternion.Euler(rotationOffset) * normalizedPoint;

    float u = 1.0f - (0.5f + Mathf.Atan2(normalizedPoint.z, normalizedPoint.x) / (2.0f * Mathf.PI) + xOffset);
    float v = 0.5f - Mathf.Asin(normalizedPoint.y) / Mathf.PI + yOffset;
    return new Vector2(u, v);
}
```

Problems

Interpolation Error



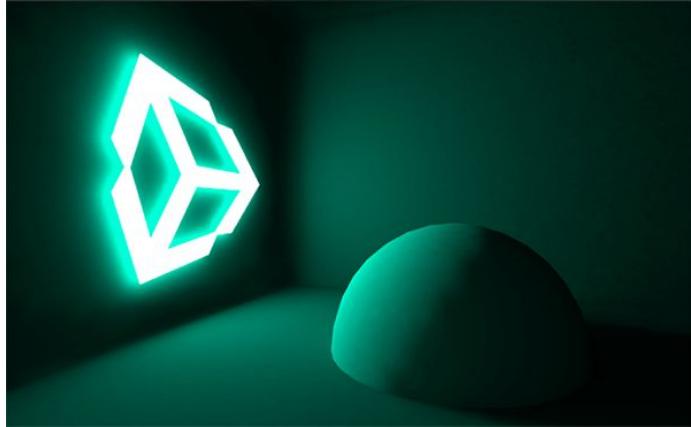
UV Seam



3. Calculating Light Probes

Realtime Global Illumination

- **Light that changes at runtime
(not baked)**
- **Direct light**
(Light Source → Object → Sensor)
- **Indirect Light**
(multiple bounces; different light sources
like emissive materials)
- **Using Enlighten middleware**



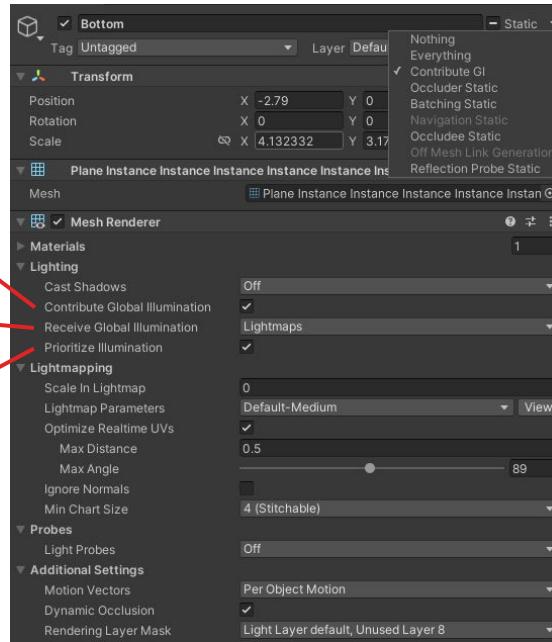
3. Calculating Light Probes using Realtime Global Illumination

Settings of the environment model

Mesh is included in GI calculations

Emission lightmap is needed to emit light

Mesh will always be included in GI calculations



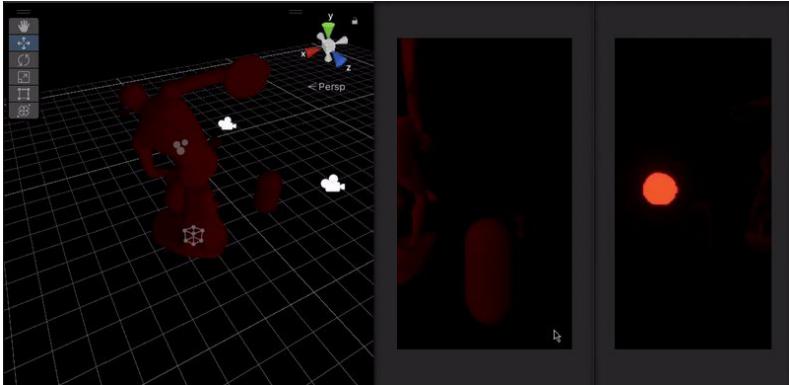
Static objects would receive lighting from other emissive objects

0 means not lightmapped but still contributes to other renderers

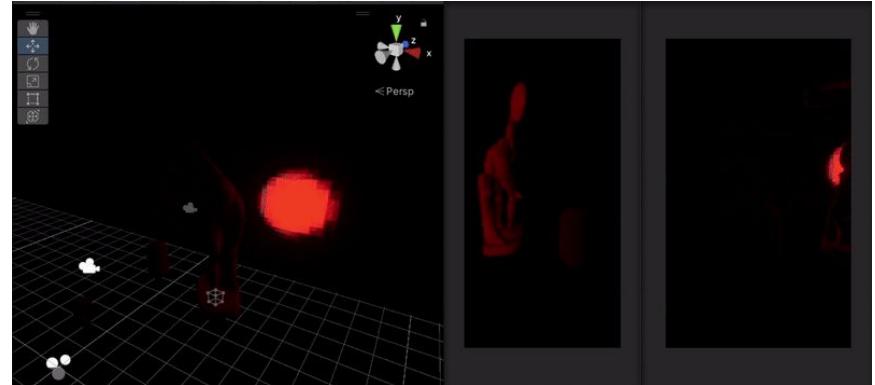
Default-Medium for better quality
Default-VeryLowResolution for higher FPS

Lightmap Parameters

Default-VeryLowResolution



Default-Medium



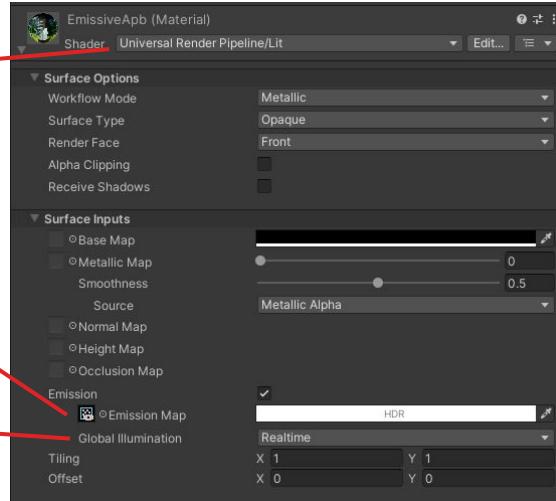
3. Calculating Light Probes using Realtime Global Illumination

Settings of the emissive material

Default shader supports
emissive settings

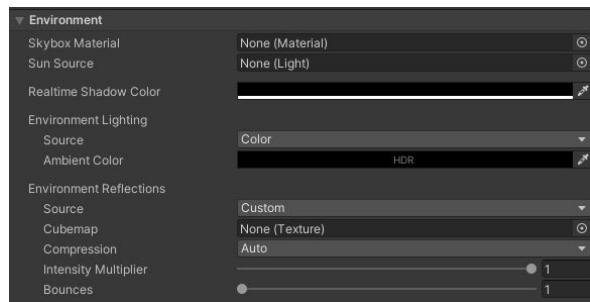
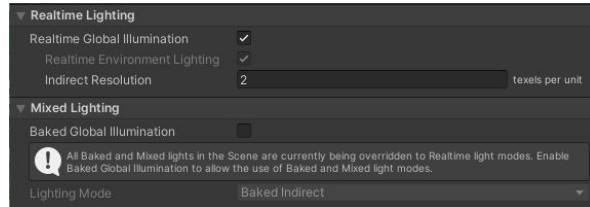
EmissionMap is set to the
webcam RenderTexture

Update GI at runtime



3. Calculating Light Probes using Realtime Global Illumination

Lighting Settings



Trigger GI Update

```
60
61     void UpdateEmissiveMaterial(GameObject go)
62     {
63         Renderer renderer = go.GetComponent<Renderer>();
64         if (renderer == null) return;
65
66         RendererExtensions.UpdateGIMaterials(renderer);
67     }
68 }
69 }
```

Light Probes need to be baked into the scene with 'Generate Lighting' button

4. Transmit Light Probes

Transport Protocol

TCP	UDP
Reliable transmission	Can be unreliable
Overhead through ACKs	No Overhead
No Broadcast	Broadcast to multiple devices

 perfect for realtime applications!

4. Transmit Light Probes

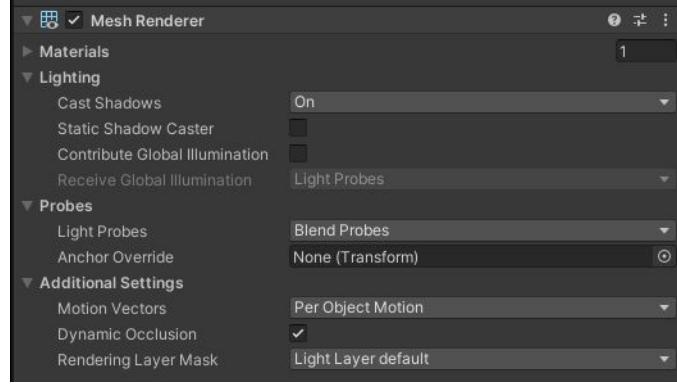
```
11
12     private UdpClient udpClient = new UdpClient();
13     public string udpIPAddress = "192.168.178.3"; // The IP address to send the data to
14     public int udpPort = 7778; // The port to send the data to
15
16     void Start()
17     {
18         udpClient.EnableBroadcast = true;
19     }
20
21     void Update()
22     {
23         SendLightProbeData(udpClient);
24     }
25
26     void SendLightProbeData(UdpClient udpClient)
27     {
28         Vector3[] probePositions = LightmapSettings.lightProbes.positions;
29         int probeCount = probePositions.Length;
30
31         // Serialize the data to a byte array
32         for (int i = 0; i < probeCount; i++)
33         {
34             SphericalHarmonicsL2 shCoefficient;
35             LightProbes.GetInterpolatedProbe(probePositions[i], null, out shCoefficient);
36             byte[] data = SerializeLightProbeData(i, shCoefficient);
37
38             // Send the data over UDP
39             udpClient.Send(data, data.Length, udpIPAddress, udpPort);
40         }
41     }
42 }
```

Disable Firewall for Unity!

	Unity 2022.3.24f1 Editor	Zulassen	Domäne
<input type="checkbox"/>	Unity 2022.3.24f1 Editor	Blockieren	Öffentlich
<input checked="" type="checkbox"/>	Unity 2022.3.37f1 Editor	Zulassen	Alle

5. Light virtual Objects

```
125     private void ReceiveData()
126     {
127         ReceiveClient = new UdpClient(udpPort);
128         ReceiveClient.EnableBroadcast = true;
129
130         while (true)
131         {
132             try
133             {
134                 IPEndPoint anyIP = new IPEndPoint(IPAddress.Any, 0);
135                 byte[] data = ReceiveClient.Receive(ref anyIP);
136
137                 DeserializeLightProbeData(data);
138             }
139             catch (Exception err)
140             {
141                 Debug.Log("<color:red>" + err.Message + "</color>");
142             }
143         }
144     }
145
146     void UpdateLightProbeSHCoefficients(SphericalHarmonicsL2[] shCoefficients)
147     {
148         LightProbes lightProbes = LightmapSettings.lightProbes;
149         SphericalHarmonicsL2[] existingBakedProbes = lightProbes.bakedProbes;
150
151         for (int i = 0; i < existingBakedProbes.Length; i++)
152         {
153             existingBakedProbes[i] = shCoefficients[i] * (1.0f / divider);
154         }
155
156         // Update the light probes in LightmapSettings
157         LightmapSettings.lightProbes.bakedProbes = existingBakedProbes;
158
159     }
160 }
```



Further Challenges

Consistent Camera Positioning

1. Place the Camera in the room
2. Guess the position in the virtual model
3. Project camera stream onto virtual model
4. Repeat from 2. until properly aligned



4 DoF

→ cumbersome process

→ take advantage of ViewR's calibration

1. Place a small virtual indicator at the position of the camera
2. Place the camera at the position of the indicator using passthrough



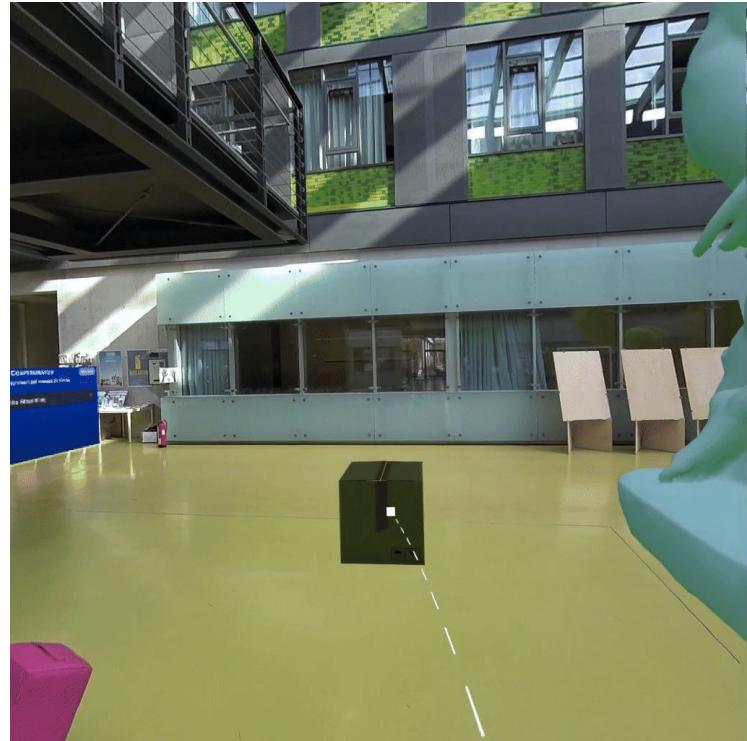
1 DoF

Camera Calibration



Further Challenges

Blue tint on the HMD



Blue tint on HMD

Initial guess was:

- Cold color temperature of the camera
- Too much influence of the blue sky



(apparently) no tint when not lit from above

Effect only visible on computer,
On HMD still blue

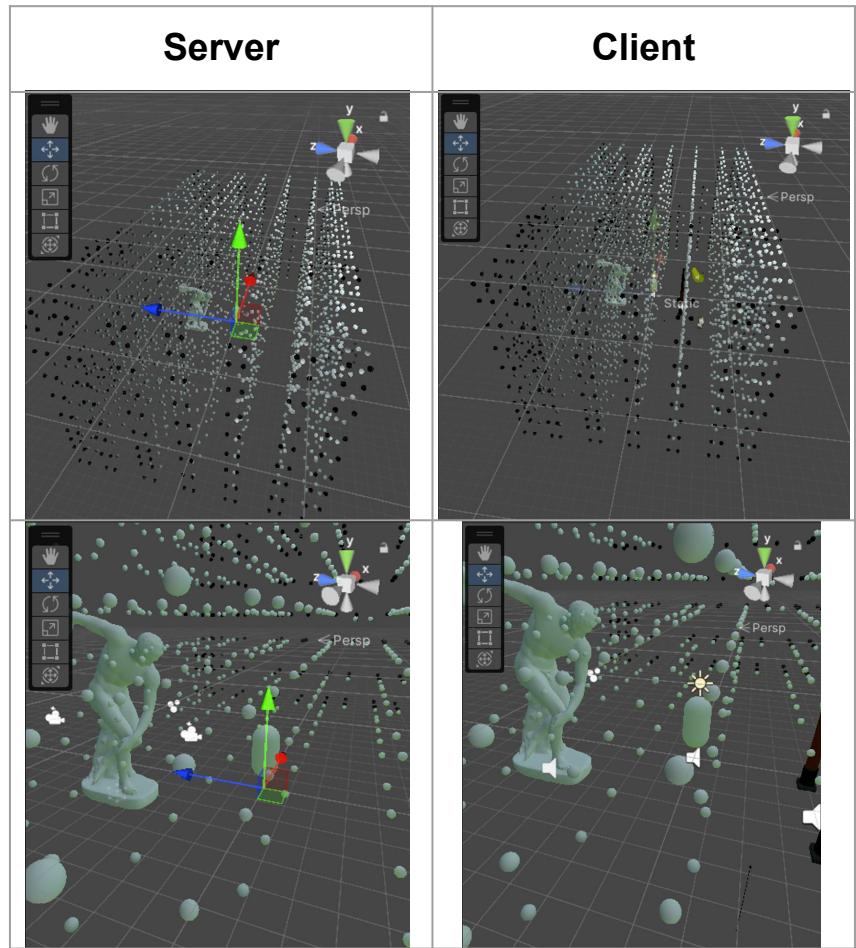
Blue tint on HMD

Next guess:

- Error in transmission
- Color channels swapped
- Wrong Byte order

→ Diff checker on all coefficients shows that transmission works

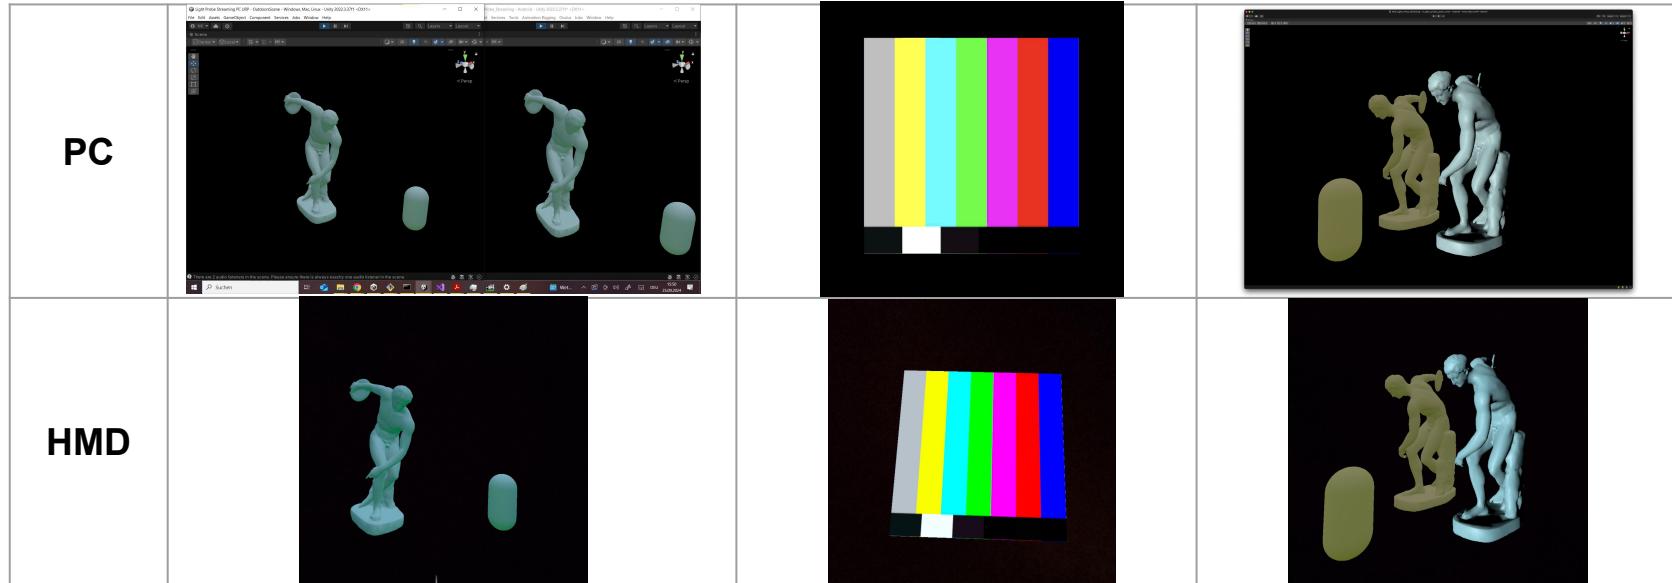
only 14 of 2359 light probes (0.6%) have discrepancies apart from rounding errors



Blue tint on HMD

Next guess:

- Display of the HMD
- Color Settings

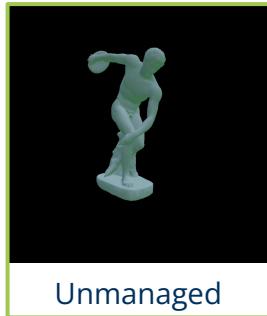


Blue tint on HMD

- Oculus Color Space Setting!



PC Reference



Unmanaged



Rift S



Rec2020



Quest 1



Rec709



DCI P3

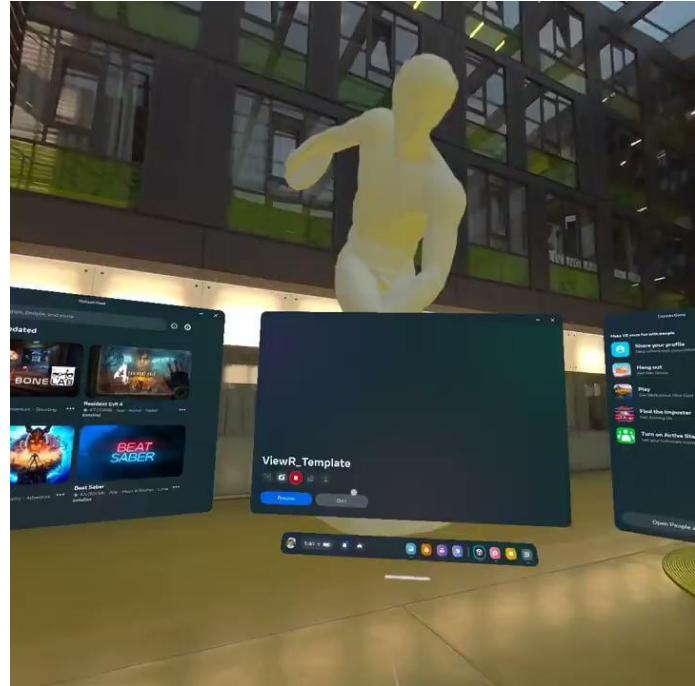


Rift CV1

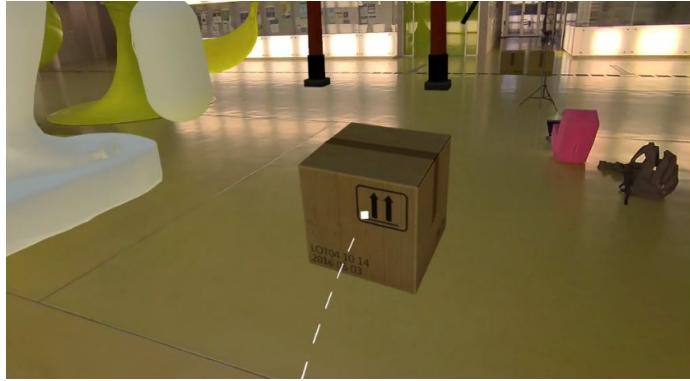


Adobe RGB

Results



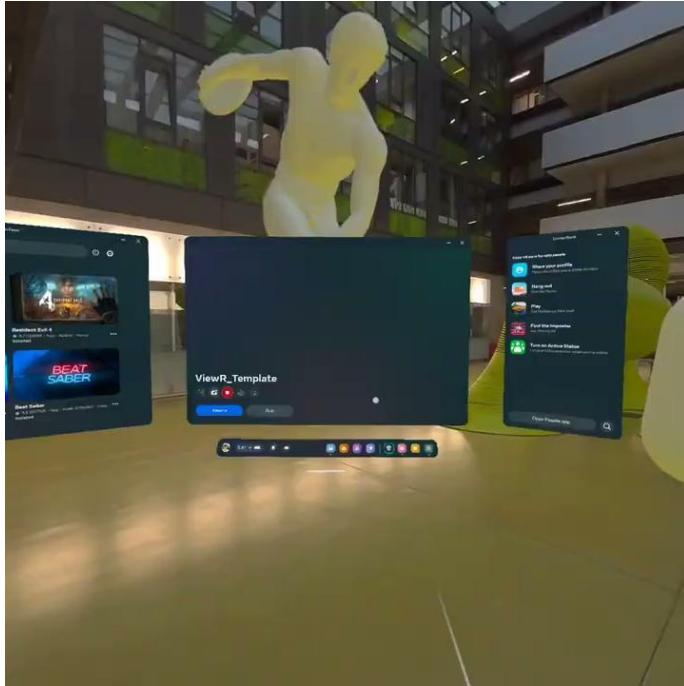
Results: Spatial Differences



Results: Spatial Differences



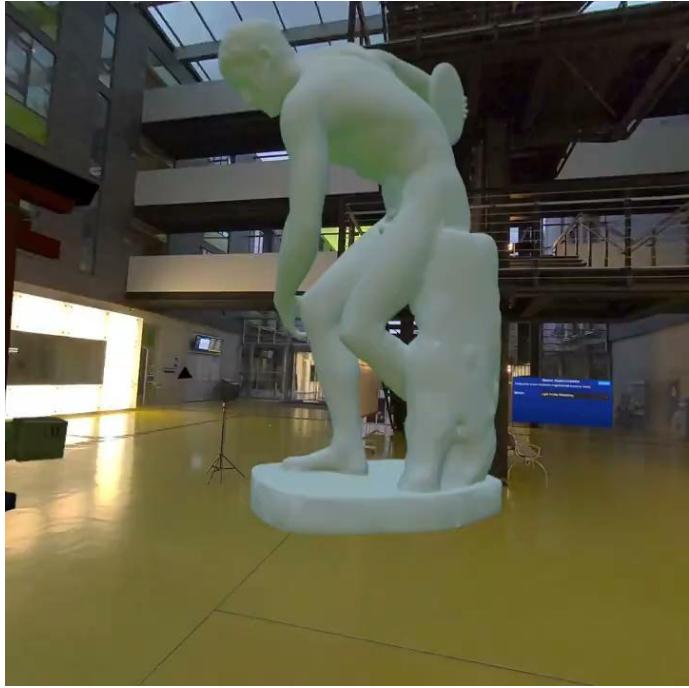
Results: Camera Occlusion



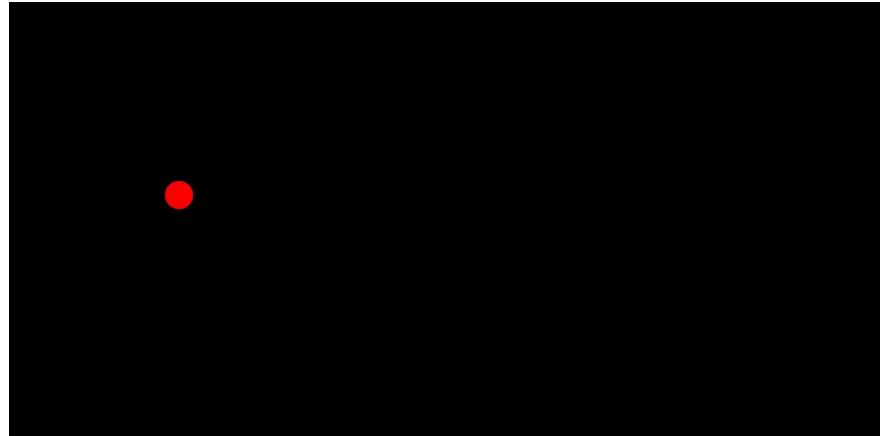
Results: Artificial Sunset



Results: Artificial Red Blobs

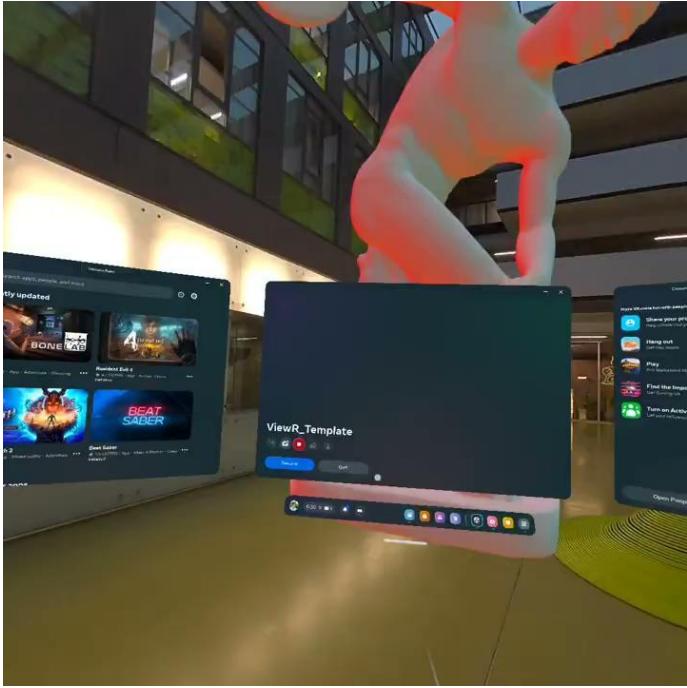


Results Artificial Red Dot



Results: Skybox comparison

Model based



Skybox



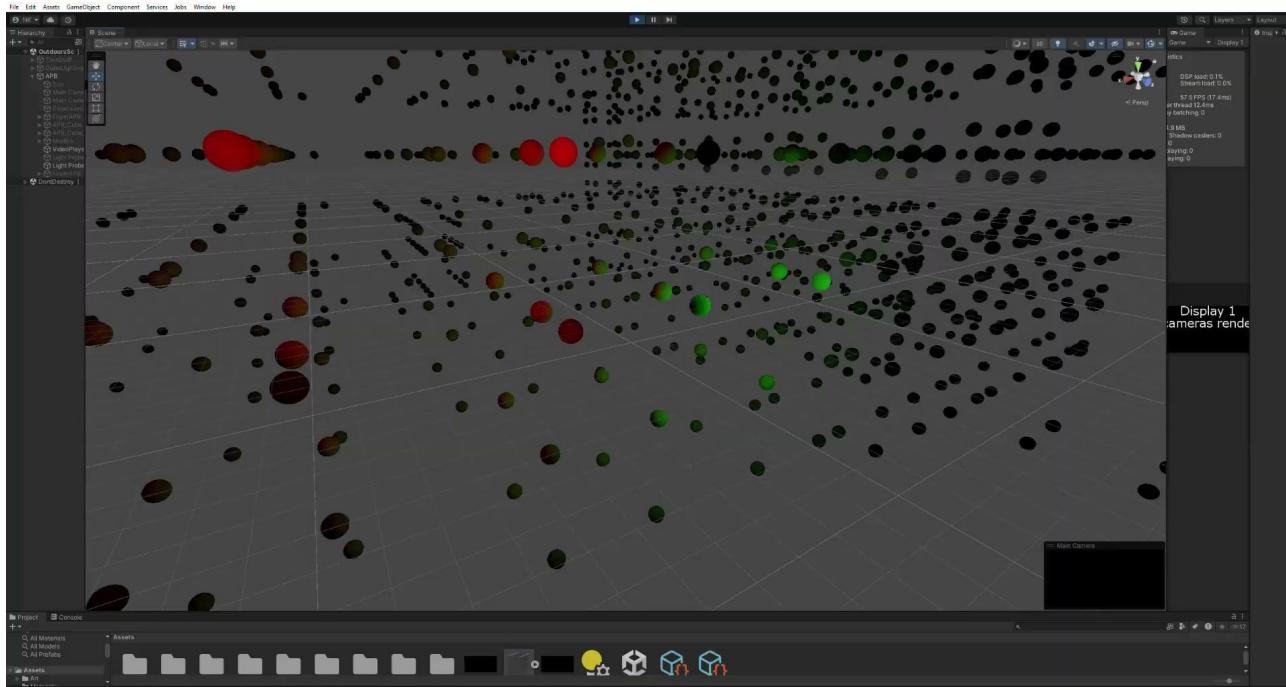
Performance

Laptop	Desktop
i5-10210U 4 Cores / 8 Threads @ 1.6GHz - 4.2GHz GeForce MX250	i7-13700K 8 P-Cores, 8 E-Cores / 16 Threads @ 2.5GHz - 5.4GHz GeForce GTX 1070
~ 10 FPS	~ 30 - 50 FPS
~ 2s delay (1s camera, 1s lighting calculation)	-

How Enlighten Realtime Global Illumination works

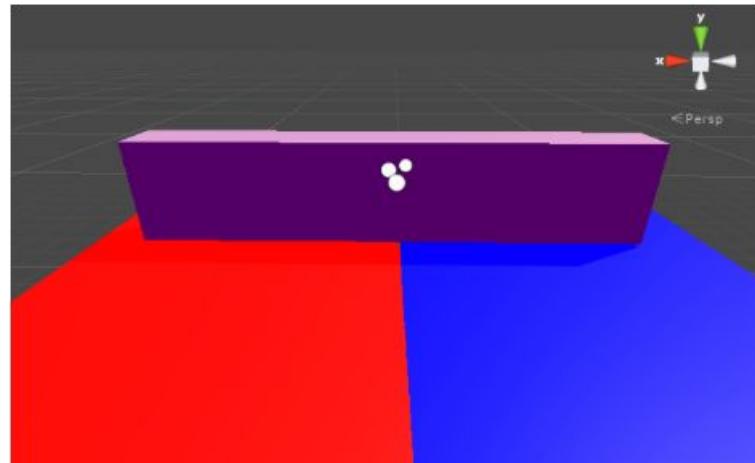
Enlighten Realtime Global Illumination splits the Scene into small surface patches and determines the degree to which these patches are visible to each other. At runtime, Enlighten Realtime Global Illumination uses this precomputed visibility information to approximate how Realtime Lights bounce in the Scene, saves the results in a set of lightmaps, and then uses these lightmaps to apply indirect lighting to the Scene. It is computationally intensive to update the lightmaps, and so the process is split across several frames. It takes Enlighten Realtime Global Illumination several frames to propagate changes to indirect lighting throughout the Scene.

Performance: Desktop PC



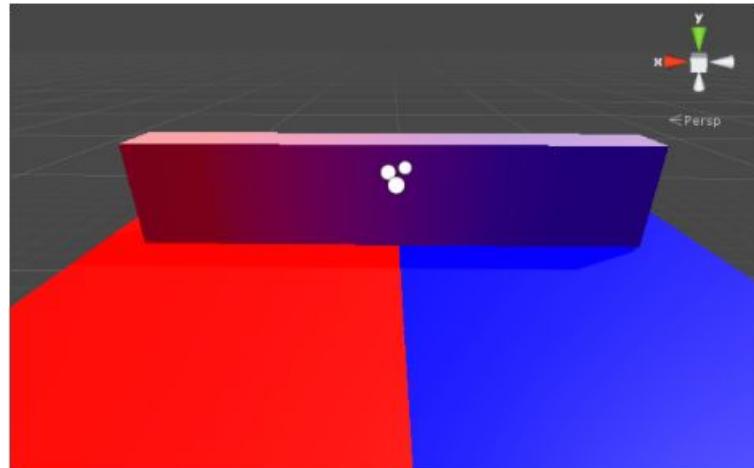
Further Considerations

Light Probe Proxy Volume



without LPPV

not supported in URP



with LPPV

Further Considerations

Adaptive Probe Volumes

*only available in
version 6 of Unity (preview)*



without APV



with APV

Further Considerations

A proper shader based projection & emission

- → less interpolation error
- → more flexible regarding repositioning of the camera

Sync interactable objects with rendering scene

- → virtual objects can interact with lighting

Place camera in the air + multiple cameras + HDR

- → reduce visual artifacts caused by humans
- → deal with obstruction
- → improve visual quality



Further Considerations

Compare to real objects



Conclusion

How can we get lighting information from the physical world to the MR application in real-time?

- **360° camera**
- **Room model + emissive material**
- **Streaming light probes over the network**

- | | |
|---|---|
| + works | - FPS & delay could be improved |
| + is visually appealing | - can't handle too dynamic environments |
| + multi-user support (tested with 2 HMDs) | - setup process is cumbersome |
| + low requirements on the HMD | |

→ works best for static setups with few concurrent users in a clear space
(e.g. exhibitions in museums)

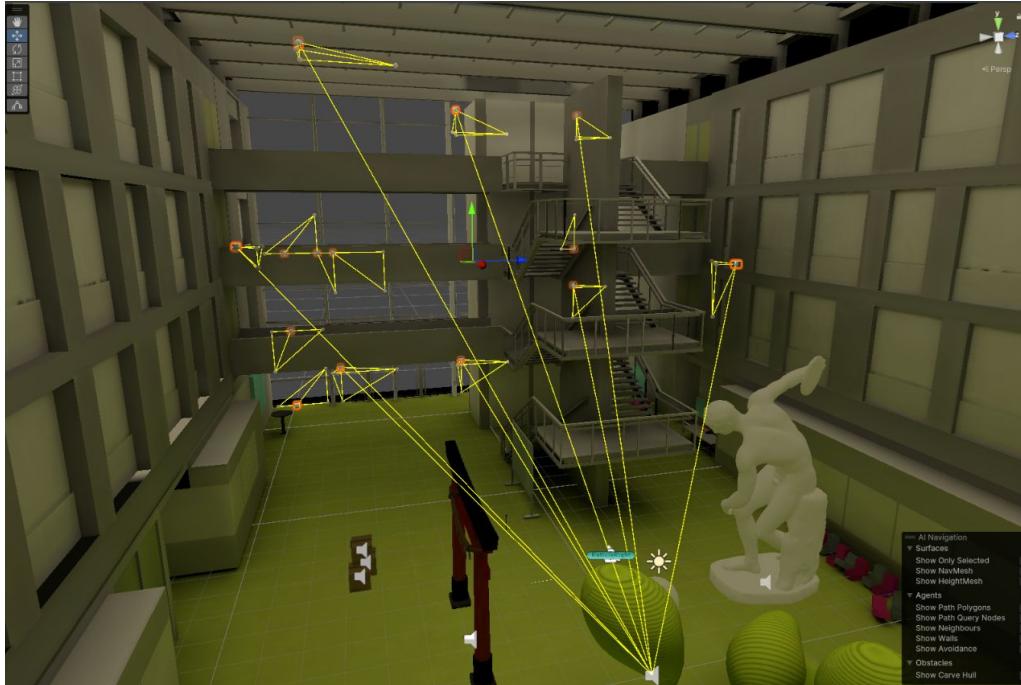
Any questions?

Image Sources

- https://scontent.oculuscdn.com/v/t64.5771-25/38974492_442085696619485_5665262356792344576_n.png?nc_cat=107&ccb=1-7&nc_s_id=6e7a0a&nc_ohc=T7N4iKmO3UQ7kNvgGuZRQ7&nc_ht=scontent.oculuscdn.com&oh=00AYBQq080owE6bwtX055LFf7dWuuPNmJXHI-IRWaM-OYsw&oe=66F8932B (slide 4)
- https://mixed-news.com/en/wp-content/uploads/2022/10/Demeo_Mixed_Reality-860x484.png (slide 4)
- <https://i.imgur.com/Fwmkpn0.png> (slide 6)
- <https://www.az-delivery.de/cdn/shop/products/tcs34725-rgb-farb-sensor-mit-infrarot-filter-diy-modul-zur-farb-erkennung-kompatibel-mit-arduino-414913.jpg?v=1679399266&width=1200> (slide 10)
- https://cdn1.botland.de/123834-pdt_540/sparkfun-tristimulus-farbsensor-opt4048dtsr-qwiic-sparkfun-sen-22638.jpg (slide 10)
- <https://de.farnell.com/productimages/standard/en/GB/3879713-40.jpg> (slide 10)
- https://m.media-amazon.com/images/I/61o2ZUzB4XL.AC_UF894,1000_QL80.jpg (slide 10)
- https://encrypted-tbn0.gstatic.com/images?q=tbn:ANd9GcRAwg2N3x8ThZlaalbuxyon9qF6mDI_pNxI_A&s (slide 15, 16)
- <https://ricohtheta.eu/cdn/shop/products/Theta-Z1-front-screen-on.jpg?v=1655989855> (slide 15, 17, 18)
- <https://vr-expert.de/wp-content/webp-express/webp-images/uploads/2023/08/Quest-3-diagnal-view.png.webp> (slide 15, 16)
- https://haid-computers.de/cdn/shop/products/rgb_gaming_pc_talius_adf23630-328c-49ea-bc12-b6bf1bf5dc38_600x.png?v=1659005283 (slide 15, 16)
- https://heise.cloudimg.io/bound/1200x1200/q85.png-lossy-85.webp-lossy-85.foil1/www-heise-de/imgs/18/3/0/8/8/9/3/9/Bildschirmfoto_2021-04-15_um_13-427ab236292eb0de.png (slide 17)
- <https://docs.unity3d.com/uploads/GlobalIllumination/EmissiveMaterial.png> (slide 22)
- <https://docs.unity3d.com/uploads/Main/LightProbeProxyVolumeExample2.png> (slide 47)
- <https://docs.unity3d.com/uploads/Main/LightProbeProxyVolumeExample1.png> (slide 47)
- <https://docs.unity3d.com/6000.0/Documentation/uploads/urp/probe-volumes/probevolumes-per-pixel.jpg> (slide 48)

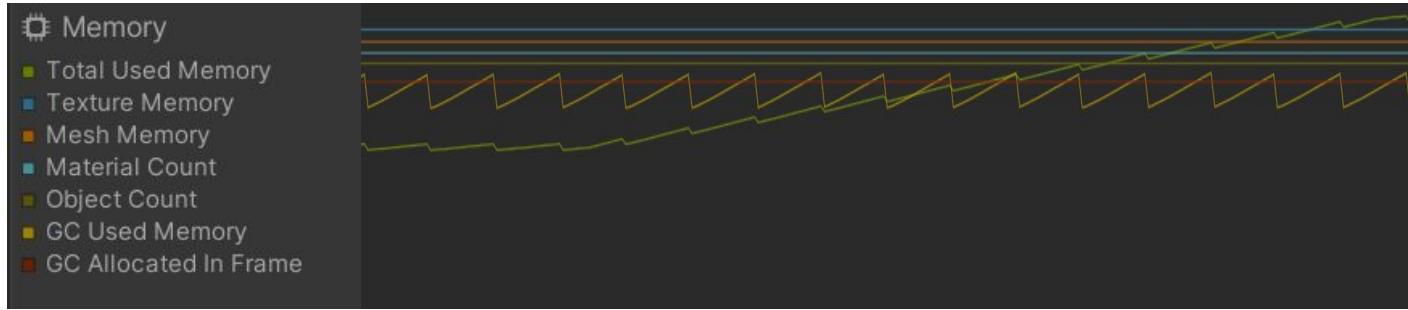
Backup Slides

14 incorrect light probes



Backup Slides

Memory Management



- SWAP used 30GB of disk space
- degraded performance of whole system

Backup Slides

Unity deleted all my Light Probes when saving the scene

- had to manually edit .unity file to get them back
- a bug in Unity?